

Mean Shift Clustering

Giacomo Ravara, Simone Casini

February 27, 2020

Università degli Studi di Firenze

giacomo.ravara@stud.unifi.it

simone.casini3@stud.unifi.it

Overview

Introduction

Sequential Version

OpenMP Version

Java Parallel Version

Results

Introduction

Mean Shift

- Mean shift is a clustering algorithm used in many applications, such as image segmentation.
- Computationally expensive
- Easy to parallelize

Mean Shift

- Non-parametric mode finding algorithm
- Define a *KDE* surface
- At each step, each point is shifted towards the nearest mode

Mean Shift

At each iteration, given a point y_i , the vector towards the next point y'_i is determined by a weighted average between y_i and each point x_j in the dataset:

$$\begin{aligned} m_i = m(y_i) &= \frac{\sum_{j=1}^n K(\|y_i - x_j\|) x_j}{\sum_{j=1}^n K(\|y_i - x_j\|)} - y_i \\ &= \frac{\sum_{j=1}^n w_{ij} x_j}{\sum_{j=1}^n w_{ij}} - y_i = \mathbf{m}_i - y_i \end{aligned}$$

The only parameters of the algorithm are the bandwidth and the choice of the *kernel* used.

- There are many different types of kernel, the most used is the *Gaussian kernel*:

$$k(x) = \frac{1}{(\sqrt{2\pi}\sigma)} e^{-\frac{x^2}{2\sigma^2}}$$

- The standard deviation σ is the bandwidth parameter, that determines the number of final clusters.

The data structure consists of two list of points:

- list containing the original points X
- list containing the points during the iterations Y

The data structures are the same for both the sequential and the parallel algorithm.

Sequential Version

Algorithm 1 Sequential Mean Shift

```
1: for each point  $X$  do
2:    $iter \leftarrow 0$ 
3:    $Y_0 \leftarrow X$ 
4:   while  $t < MaxIterations$  do
5:      $Y_{t+1} = \text{SHIFTER}(Y_t, X, bandwidth)$ 
6:      $\text{SWAP}(Y_t, Y_{t+1})$ 
7:      $t \leftarrow t + 1$ 
8:   Return  $Y_{t+1}$ 
```

OpenMP Version

- OpenMP is an API for shared-memory programming.
- The program is parallelized with *compiler directives*
- It follows a *fork-join* thread model.

Mean Shift OpenMP

- Mean Shift is easily parallelizable.
- Each point is shifted independently from the other points.
- In OpenMP we use the directive *pragma* to parallelize.

Algorithm 2 OpenMP Mean Shift

```
1: #pragma
2: for each point  $X$  do
3:    $iter \leftarrow 0$ 
4:    $Y_0 \leftarrow X$ 
5:   while  $t < MaxIterations$  do
6:      $Y_{t+1} = \text{SHIFTER}(Y_t, X, bandwidth)$ 
7:      $\text{SWAP}(Y_t, Y_{t+1})$ 
8:      $t \leftarrow t + 1$ 
9:   Return  $Y_{t+1}$ 
```

Implementation

The best scheduling strategy for the mean shift algorithm is the *dynamic scheduling*. Using this strategy ensure that each thread will never stops its execution waiting for other threads.

Implementation

Threads	Static	Dynamic
2	124.54s	114.49s
3	87.07s	76.29s
4	63.72s	57.20s
5	52.44s	45.72s
6	44.13s	38.25s

Comparison between static and dynamic with 36864 points

Java Parallel Version

Java Parallel Version

- Uses the same data structures.
- Divide points between threads.
- Each point is shifted until convergence.

Algorithm 3 Java Mean Shift

- 1: **for** $t < num_threads$ **do**
- 2: `CREATETHREADMEANSHIFT` ($Y_t, X, bandwidth$)
- 3: `SHIFTER`

Return ShiftedPoints

Implementation

It is possible an alternative implementation, without a static division of points between threads.

The threads share an index to iterate the shared list of points, the index is defined as an `AtomicInteger` to avoid race condition.

Implementation

Threads	Static	Dynamic
2	135.10s	133.20s
3	93.90s	88.35s
4	77.87s	66.68s
5	64.25s	53.27s
6	56.26s	45.09s

Comparison between static and dynamic with 36864 points

Results

- The performances are measured by evaluating the speedup
- Changing the number of threads and using different input sizes
- Using Microsoft Azure Virtual Machine, 2.60 GHz with 6 cores and NVIDIA Tesla M60.

Results

BandWidth	1	
Number of points	16384	
Sequential C++	25.05s	
Sequential Java	25.24s	
Threads	OpenMP	Java
2	12.54s (1.99X)	13.07s (1.93X)
3	8.38s (2.98X)	8.62s (2.93X)
4	6.31s (3.96X)	6.52s (3.86X)
5	5.05s (4.95X)	5.25s (4.80X)
6	4.24s (5.90X)	4.47s (5.64X)

Tabelle 1: Comparison between sequential and OpenMP/Java

Results

BandWidth	10	
Number of points	16384	
Sequential C++	31.74s	
Sequential Java	33.34s	
Threads	OpenMP	Java
2	15.07s (2.10X)	17.62s (1.89X)
3	10.05s (3.15X)	11.70s (2.84X)
4	7.53s (4.21X)	8.83s (3.77X)
5	6.04s (5.24X)	7.11s (4.68X)
6	5.04s (6.29X)	5.99s (5.56X)

Results

BandWidth	25	
Number of points	16384	
Sequential C++	47.01s	
Sequential Java	50.30s	
Threads	OpenMP	Java
2	23.18s (2.03X)	25.77s (1.95X)
3	15.52s (3.03X)	17.18s (2.92X)
4	11.59s (4.06X)	12.89s (3.90X)
5	9.26s (5.08X)	10.30s (4.84X)
6	7.76s (6.06X)	8.74s (5.75X)

Results

BandWidth	25		
Number of points	4096		
Sequential C++	2.86s		
Sequential Java	3.05s		
Threads	OpenMP	Java	
2	1.43s (2.00X)	1.64s (1.80X)	
3	0.98s (2.92X)	1.03s (2.94X)	
4	0.72s (3.94X)	0.78s (3.88X)	
5	0.57s (4.97X)	0.64s (4.74X)	
6	0.47s (6.03X)	0.57s (5.30X)	

Results

BandWidth	25		
Number of points	36864		
Sequential C++	228.61s		
Sequential Java	245.05s		
Threads	OpenMP	Java	
2	114.37s (1.99X)	131.75s (1.86X)	
3	76.30s (2.99X)	87.96s (2.79X)	
4	57.05s (4.00X)	66.02s (3.72X)	
5	45.77s (4.99X)	53.46s (4.59X)	
6	38.20s (5.98X)	44.88s (5.46X)	

Results



Example of Image Segmentation, using MeanShift

END